

CODED - Cloud-Optimized Distributed Earth Data

Can the decentralized peer-to-peer **InterPlanetary File System (IPFS)** make important environmental datasets resilient against institutional takedowns, while remaining useful for cloud-native geospatial workflows? (Zarr · Icechunk · xarray)

Cory Levinson · Rich Signell · Brian Terry

★ An ESIP Funding Friday Project

Objective

The geoscience community has converged on **cloud-optimized formats** - Zarr, and now **Icechunk** - that stream huge datasets via chunk-level range requests. **IPFS** offers content-addressed, location-independent, takedown-resistant storage. Can the two be married?

This round we go past "does it work" to "where does the time go" - separating discovery, transfer, and caching costs.

Preserve your dataset - CID on Filebase

Have an existing **Zarr** or **Icechunk** store you can't afford to lose? Add it to IPFS for a content ID (CID), then pin that CID on **Filebase** (IPFS + Filecoin, free to 5 GB) - or **any public IPFS node** - so it stays available, verifiable, and independent of any single institution.

```
# 1. add your existing store to IPFS → root CID
cid=$(ipfs add -rQ --cid-version=1 my.icechunk/)
# 2. export the whole DAG (all chunks + refs) to a CAR file
ipfs dag export $cid > my.car
# 3. pin that CAR on Filebase – import=car preserves the CID
aws s3 cp my.car s3://bucket/my.car \
  --endpoint-url https://s3.filebase.com \
  --metadata import=car # same CID, now persistent
```

The **CID is a hash of the data** — tamper-evident and location-independent. Share it beside your DOI.

Icechunk reads IPFS - no adapter


Icechunk 2.0.5 ships `http_storage(base_url=...)`, a read-only store over HTTP. Point it at an IPFS gateway path and the entire repo opens in xarray - the on-disk layout survives `ipfs add -r` untouched.

```
# whole Icechunk repo, straight from a CID
st = icechunk.http_storage("https://gw/ipfs/<CID>")
repo = icechunk.Repository.open(st)
sess = repo.readonly_session("main")
ds = xr.open_zarr(sess.store)
```

Old CID → old snapshot (reproducibility). Edit one of four chunks → **6 shared blocks reused** — cross-version dedup for free.

Caching is the real win

Once a block has been fetched, the second read serves from the local blockstore - no network round-trip. Cold→warm gains were dramatic on the long link.



Warm reads converge to ~11 s / ~190 MB/s regardless of topology. A local pin **is** the warm floor.

Where does the time go?

We read the **same 2.08 GB ERA5** 2-m temperature slice (500 Zarr chunks) three ways, in two topologies, splitting each read into **discovery** vs **transfer**. Every path returns a bit-identical mean = **286.5062 K**.

Cold read path	Same-region	Cross-region
Local pin (warm floor)	~11 s	~11 s
Direct peer (Bitswap)	+16 s	+57 s
DHT discovery vs direct	~noise	~noise
Icechunk on S3 (control)	51 s	51 s

Transfer overhead = direct - local. Same-region ≈130 MB/s VPC; cross-region ≈36 MB/s WAN. Discovery (DHT vs direct) is within noise in a 2-node swarm - a lower-bound caveat, not a general claim.

What This Means

- **Blame the distance, not the DHT.** Discovery is cheap; high-RTT single-peer transfer is the cost. More peers & regional gateways close the gap.
- **Warm ≈ free.** Caching / pinning gives repeat readers a local-speed floor - the natural pattern for shared community data.
- **Icechunk-on-IPFS is real today** via `http_storage` - reproducible, dedup-friendly, zero adapter code.

Bottom Line
Distance dominates - not discovery. Finding content on IPFS is cheap; moving 2 GB across a continent is the cost, and the second read is nearly free. Co-locate a gateway (or pin locally) and IPFS delivers S3-class throughput with verifiable, takedown-resistant data.

Acknowledgements: This work was greatly assisted by an OpenClaw agent running on an AWS VM using Claude Sonnet through Amazon Bedrock. Thanks to AWS for providing ESIP Credits to do this work!