

CODED - Cloud-Optimized Distributed Earth Data

Can the decentralized peer-to-peer **InterPlanetary File System (IPFS)** make important environmental datasets resilient against institutional takedowns, while remaining useful for cloud-native geospatial workflows? (Zarr · Icechunk · xarray)

Cory Levinson · Rich Signell · Brian Terry

★ An ESIP Funding Friday Project

Objective

The geoscience community has converged on **cloud-optimized formats** - Zarr, and now **Icechunk** - that stream huge datasets via chunk-level range requests. **IPFS** offers content-addressed, location-independent, takedown-resistant storage. Can the two be married?

This round we go past "does it work" to "where does the time go" - separating discovery, transfer, and caching costs.

Preserve *your* dataset - CID on Filebase

Have an existing **Zarr** or **Icechunk** store you can't afford to lose? Add it to IPFS for a content ID (CID), then pin that CID on **Filebase** (IPFS + Filecoin, free to 5 GB / 500 pins) - or **any public IPFS node** - so it stays available, verifiable, and independent of any single institution.

```
# 1. add your existing store to IPFS → root CID
cid=$(ipfs add -rQ --cid-version=1 my.icechunk/)
# 2. export the whole DAG (all chunks + refs) to a CAR file
ipfs dag export $cid > my.car
# 3. pin that CAR on Filebase — import=car preserves the CID
aws s3 cp my.car s3://bucket/my.car \
  --endpoint-url https://s3.filebase.com \
  --metadata import=car # same CID, now persistent
```

The **CID is a hash of the data** — tamper-evident and location-independent. Our live ERA5 store:

bafybeiecltl3mtags2i3jeumxe5c7zuvj2r76zxwsg6tfljiouvvywqzbaq

One CID = the whole queryable store. Re-chunked to 100 blocks (5 timesteps each) to fit the free-tier pin budget - **data byte-identical** (maxdiff 0.0).

Icechunk reads IPFS - no adapter


Icechunk 2.0.5 ships `http_storage(base_url=...)`, a read-only store over HTTP. Point it at an IPFS gateway path and the entire repo opens in xarray - the on-disk layout survives `ipfs add -r` untouched.

```
# whole Icechunk repo, straight from a CID
st = icechunk.http_storage("https://gw/ipfs/<CID>")
repo = icechunk.Repository.open(st)
sess = repo.readonly_session("main")
ds = xr.open_zarr(sess.store)
```

Old CID → old snapshot (reproducibility). Edit one of four chunks → **6 shared blocks reused** — cross-version dedup for free.

Caching is the real win

Once a block has been fetched, the second read serves from the local blockstore - no network round-trip. Cold→warm gains were dramatic on the long link.



Warm reads converge to ~11 s / ~190 MB/s regardless of topology. A local pin **is** the warm floor.

Where does the time go?

We read the **same 2.08 GB ERA5** 2-m temperature slice (500 Zarr chunks) three ways, in two topologies, splitting each read into **discovery** vs **transfer**. Every path returns a bit-identical mean = **286.5062 K**.

Cold read path	Same-region	Cross-region
Local pin (warm floor)	~11 s	~11 s
Direct peer (Bitswap)	+16 s	+57 s
DHT discovery vs direct	~noise	~noise
Icechunk on S3 (control)	51 s	51 s

Transfer overhead = direct - local. Same-region ≈130 MB/s VPC; cross-region ≈36 MB/s WAN. Discovery (DHT vs direct) is within noise in a 2-node swarm - a lower-bound caveat, not a general claim.

What This Means

- **Blame the distance, not the DHT.** Discovery is cheap; high-RTT single-peer transfer is the cost. More peers & regional gateways close the gap.
- **Warm ≈ free.** Caching / pinning gives repeat readers a local-speed floor - the natural pattern for shared community data.
- **Icechunk-on-IPFS is real today** via `http_storage` - reproducible, dedup-friendly, zero adapter code.

Bottom Line
Distance dominates - not discovery. Finding content on IPFS is cheap; moving 2 GB across a continent is the cost, and the second read is nearly free. Co-locate a gateway (or pin locally) and IPFS delivers S3-class throughput with verifiable, takedown-resistant data.

